

Lecture 15: The transformer model

PSYC 51.17: Models of language and communication

Jeremy R. Manning
Dartmouth College
Winter 2026

Learning objectives

By the end of this lecture, you will be able to...

1. Understand the transformer as a function that predicts the next word
2. Follow data from raw text through tokenization, embedding, and attention
3. Explain how queries, keys, and values enable self-attention
4. See how multiple attention heads capture different patterns
5. Understand how stacking blocks creates deep representations

The transformer model

Further reading

See [The Animated Transformer](#) for a visual walkthrough of the transformer architecture. (Today's lecture is heavily inspired by this resource!)

Definition

The transformer is a machine learning model for **sequence modeling**. Given a sequence of *things*, the model can predict what the next *thing* in the sequence might be.

The big picture

We can think of the transformer as a function: $\text{Transformer}(X, \theta) \rightarrow Y$

- X is our input sequence
- θ represents the model parameters
- Y is the predicted next token

The transformer model

Tokenization

Why tokenize?

The transformer operates on sequences of *numbers*. So first we need to **tokenize** the input text into discrete units (tokens) and map those tokens to unique IDs. The model doesn't understand tokens directly—it identifies them using unique numbers from a vocabulary.

Our running example

"the robots will bring" → [3206, 2736, 3657, 400]

Each word maps to a unique ID in the vocabulary. (In practice, tokenization is more complex—subword units, punctuation, etc.—but this simplified example suffices for our purposes.)

Embeddings: numbers speak louder than words

Remember embeddings?

The transformer maintains **embeddings** of each token in its vocabulary. These embeddings are learned vectors that capture semantic information about the tokens.

Recall from last week...

This is the same idea as LSA, LDA, Word2Vec, GloVe, FastText, etc.- but here the embeddings are learned jointly with the rest of the model.

Token embeddings

Token embeddings

Dimensions

Our transformer has embedding vectors of length $C = 768$. All embeddings can be packed together in a single $T \times C$ matrix, where (in our example) $T = 4$ is the number of input tokens.

Intuition

Each *row* of this matrix corresponds to a token in the input sequence, and each *column* corresponds to a dimension in the embedding space.

Position embeddings

Position embeddings

Constructing position embeddings

Each position in the input sequence (1st, 2nd, 3rd, etc.) also has a learned embedding vector of length $C = 768$. These position embeddings are stored in a separate $T \times C$ matrix.

Why positions matter

In order to capture the significance of the **position** of a token within a sequence, the transformer also maintains embeddings for each position.

Without position information, "cat sat" and "sat cat" would look identical to the model!

Each *row* of this matrix corresponds to a position in the input sequence, and each *column* corresponds to a dimension in the embedding space.

Token vs. position embeddings

Notice that both token and position embeddings have the same shape: $T \times C$. This allows the model to capture both the identity of the token and its position in the sequence within compatible embedding spaces.

Combined embeddings

Token embedding matrix

$T \times C$

$$\begin{matrix} \text{the} \\ \text{robots} \\ \text{will} \\ \text{bring} \end{matrix} \begin{bmatrix} 0.07 & 0.13 & 0.63 & \dots & 0.23 \\ 0.81 & 0.51 & 0.44 & \dots & 0.98 \\ 0.62 & 0.29 & 0.80 & \dots & 0.34 \\ 0.50 & 0.16 & 0.93 & \dots & 0.19 \end{bmatrix}$$

+

Position embedding matrix

$T \times C$

$$\begin{matrix} 0 \\ 1 \\ 2 \\ 3 \end{matrix} \begin{bmatrix} 0.32 & 0.12 & 0.88 & \dots & 0.75 \\ 0.14 & 0.07 & 0.41 & \dots & 0.35 \\ 0.07 & 0.47 & 0.23 & \dots & 0.86 \\ 0.81 & 0.12 & 0.06 & \dots & 0.24 \end{bmatrix}$$

Combined embeddings

Token + position = input

The token and position embedding matrices ($T \times C$ each) are **added together** to obtain a position-dependent embedding for each token.

The token and position embeddings are all part of θ , meaning they are tuned during model training.

Why is this useful?

Adding the embeddings allows the model to consider both the identity of the token and its position in the sequence simultaneously.

Queries, keys, and values

Self-attention

The Transformer computes three vectors for each token: **query**, **key**, and **value**. This is done by multiplying with learned weight matrices:

$$Q = XW_Q \quad K = XW_K \quad V = XW_V$$

The weight matrices W_Q , W_K , W_V are all part of θ .

Dimensions

The W_Q , W_K , and W_V matrices each have shape $C \times C$. Multiplying the input matrix X ($T \times C$) by these weight matrices results in three new $T \times C$ matrices: Q , K , and V .

Queries, keys, and values

$$\begin{array}{l} \text{the} \\ \text{robots} \\ \text{will} \\ \text{bring} \end{array} \begin{array}{c} T \times C \\ \left[\begin{array}{cccc} 0.39 & 0.25 & \dots & 0.98 \\ 0.95 & 0.58 & \dots & 1.33 \\ 0.69 & 0.77 & \dots & 1.20 \\ 1.31 & 0.28 & \dots & 0.43 \end{array} \right] \times \begin{array}{c} W_q (C \times C) \\ \left[\begin{array}{cccc} 0.02 & 0.21 & \dots & 0.37 \\ 0.03 & 0.02 & \dots & 0.05 \\ 0.29 & 0.07 & \dots & 0.99 \\ 0.38 & 0.37 & \dots & 0.28 \\ \vdots & \vdots & \vdots & \vdots \end{array} \right] = \end{array}$$

Queries, keys, and values

First let's start with an analogy

Imagine you have a database of images with text descriptions:

- **Query:** the user's search text
- **Key:** the text descriptions in your database
- **Value:** the actual images

Only those images (values) whose descriptions (keys) best match the search (query) are returned.

Self-attention intuition

Self-attention works similarly—tokens "query" other tokens to find which ones they should "pay attention" to. In this context, "paying attention" means incorporating information from relevant tokens into their own representation (i.e., their row in the output matrix).

Many heads are better than one

Multiple attention heads

The Transformer splits the Q , K , V matrices into multiple **heads**. With $C = 768$ columns and 12 heads, each head operates on 64 dimensions.

Why multiple heads?

Different heads can specialize in different patterns: syntax, co-reference, semantic roles, etc. The model "decides" (i.e., learns from training data) what's useful!

Splitting into heads

$Q (T \times C)$

$$\begin{bmatrix} 0.07 & 0.04 & \dots & 0.64 \\ 0.07 & 0.07 & \dots & 0.36 \\ 0.08 & 0.06 & \dots & 0.44 \\ 0.04 & 0.04 & \dots & 0.34 \end{bmatrix}$$

Time to pay attention

Self-attention

Self-attention is the core idea behind the transformer. We compute an **attention scores** matrix:

$$A = \frac{Q \cdot K^T}{\sqrt{H}}$$

This matrix tells us how much attention each token should pay to every other token.

Dimensions

- Q is $T \times C$
- K^T is $C \times T$
- A is $T \times T$
- H is the head dimension (e.g., 64)

Two things to notice

1. A can be enormous for long sequences (e.g., 2048 tokens $\rightarrow$ 2048×2048 matrix). This is a key computational bottleneck.
2. Notice that the " V " matrix isn't used yet; it'll come back soon!

Computing attention scores

$$\begin{array}{l} \text{the} \\ \text{robots} \\ \text{will} \\ \text{bring} \end{array} \begin{array}{c} Q_1 (T \times H) \\ \left[\begin{array}{cccc} 0.07 & 0.04 & \dots & 0.64 \\ 0.07 & 0.07 & \dots & 0.36 \\ 0.08 & 0.06 & \dots & 0.44 \\ 0.04 & 0.04 & \dots & 0.34 \end{array} \right] \end{array} \times \begin{array}{c} K_1 (T \times H) \\ \left[\begin{array}{cccc} 0.22 & 0.13 & \dots & 0.10 \\ 0.22 & 0.03 & \dots & 0.10 \\ 0.20 & 0.04 & \dots & 0.11 \\ 0.19 & 0.06 & \dots & 0.24 \end{array} \right] \end{array} \begin{array}{l} \text{the} \\ \text{robots} \\ \text{will} \\ \text{bring} \end{array}$$

Applying attention

Causal masking

The attention score for a token needs to be **masked** if it occurs later in the sequence. In our example, "bring" can pay attention to "robots", but not vice-versa—a token shouldn't look at future tokens when predicting its next token.

Bringing back V

To get the final output, we need to apply the attention scores to the value matrix V by multiplying A and V .

Intuition

V contains the information we want to aggregate, while A tells us *how much* of each token's value to include in the output for each token. We end up with a new $T \times C$ matrix where each token's representation is a weighted sum of the value vectors of all tokens.

Implementation details

1. **Mask** future tokens (set upper triangle to $-\infty$)
2. **Softmax** each row (normalize to probabilities)
3. **Multiply by V** (weighted sum of value vectors)

Softmax definition

The softmax function converts raw scores into probabilities (note: $-\infty$ becomes 0 after softmax):

$$\text{softmax}(z_i) = \frac{e^{z_i}}{\sum_j e^{z_j}}$$

Applying attention

$A_1 (T \times T)$

the	0.18	0.22	0.23	0.37
robots	0.18	0.29	0.30	0.24
will	0.18	0.26	0.28	0.28
bring	0.17	0.30	0.31	0.22
	the	robots	will	bring

Merging heads

Combining heads

After computing attention outputs for each head, we **concatenate** the outputs from all heads along the feature dimension to form a single $T \times C$ matrix.

Why concatenate?

Each head captures different aspects of the input sequence. By concatenating them, we preserve the diverse information learned by each head.

Intuition 1

Remember how in ELIZA our text processing function computed the *output* (response) by combining various features of the input? The effect of applying attention and merging heads is similar: we get a rich representation of the input sequence that informs us about *future* tokens. This will be used in the next step to make predictions about the *next* token in the sequence.

Intuition 2

The effect of adding in positional information and applying attention is that each token's representation now encodes not just its identity, but also its context within the sequence. This "sneaks in" more learnable parameters that enable the model to capture rich temporal structure in the data.

Merging heads

Feed forward

Feed-forward network (FFN)

A **feed-forward network** processes each token's representation independently using two linear transformations with a non-linear activation in between:

$$\text{FFN}(x) = \text{ReLU}(xW_{nn} + b_1)W_{mm} + b_2$$

The hidden layer expands to $4C = 3072$ dimensions, then projects back to $C = 768$.

Dimensions

- W_{nn} : $C \times 4C$
- b_1 : $4C$ (gets replicated for each token; not shown in next animation)
- W_{mm} : $4C \times C$
- b_2 : C (gets replicated for each token; not shown in next animation)

Why is this important?

Everything so far has used **linear operations**. The non-linearity (ReLU) in the FFN allows the model to learn complex, non-linear relationships.

Feed-forward network

Feed-forward network

Why it matters

All the weight matrices in the FFN are part of θ (i.e., the model parameters) and are learned during training. The FFN allows the model to transform the attention outputs into richer representations before passing them to the next layer.

Intuition

The expansion to $4C$ dimensions allows the model to capture more complex patterns, while the projection back to C ensures that the output remains compatible with the rest of the model.

We need to go deeper


Transformer block

All the steps thus constitute a *single transformer block*. Each block takes a $T \times C$ matrix as input and outputs a $T \times C$ matrix.

Model sizes

- GPT-2: 12–48 blocks
- GPT-3: 96 blocks
- GPT-4+: 100+ blocks (exact architecture not public)

Stacking transformer blocks



Making a prediction

The final step

Take the last token's output vector ($1 \times C$) and multiply by a $C \times V$ weight matrix, where V is the vocabulary size. Normalize (softmax) to get a probability distribution over all words.

How does it work?

In our example, "prosperity" might get a 92% probability, but "suffering" only 20%, and so on. So if "prosperity" has the highest probability, the model predicts "the robots will bring **prosperity**"

Why use only the last token?

We only care about predicting the *next* token after the entire input sequence, so we only use the output corresponding to the last token. But we needed to process the entire sequence to get context!

Making a prediction

Text generator go brrr

Max prompt length = 5

the robots will bring

Transformer

Autoregressive generation

How text generation works

The first token produced is added to the prompt and fed back to produce the second token, which is then fed back to produce the third, and so on.

Context limit

Transformers have a maximum context length (N tokens). As generation continues, we drop the oldest tokens.

Impact of context length

Remember that A matrix? Its size grows quadratically with context length ($N \times N$). Longer contexts require more memory and computation, which can be a bottleneck. This limits how far back the model can "remember" during generation: **tokens prior to the start of the context window are completely invisible to the model!**

And...that's it!

The complete picture

1. **Tokenize** text to IDs
2. **Embed** tokens + positions
3. **Project** to Q, K, V
4. **Split** into attention heads
5. **Compute** attention scores
6. **Apply** masking and softmax
7. **Multiply** by V for output
8. **Feed forward** with non-linearity
9. **Stack** many blocks
10. **Predict** next token

...Mostly!

Some things we skipped

To focus on the most important aspects, we skipped:

- **Layer normalization** (stabilizes training)
- **Residual connections** (helps gradient flow)
- **Dropout** (regularization)
- **Training** (how θ is learned)

See [nanoGPT](#) for a complete implementation.

References

Further reading

Vaswani et al. (2017, *arXiv*) "Attention Is All You Need" — The original transformer paper.

The Animated Transformer — Visual walkthrough that inspired this lecture.

The Illustrated Transformer — Jay Alammar's detailed visual guide.

nanoGPT — Andrej Karpathy's minimal GPT implementation.


Questions?



Email



Discord



Office hours

Up next...

Training transformers: where do the parameters (θ) come from?