

@keybeamerframestandout[true]

@reseteecodes

Week 6, Lecture 3 - From Theory to Practice

PSYC 51.07: Models of Language and Communication

Winter 2026

Today's Agenda

-  **Real-World Applications:** Where BERT shines
-  **Cognitive Neuroscience:** Brain-model parallels
-  **Understanding vs. Pattern Matching:** The big debate
-  **Limitations:** What BERT can't do
-  **Practical Tips:** Deployment and optimization
-  **Future Directions:** Where are we heading?

Goal: Connect BERT to real applications and understand broader implications

BERT Applications

BERT excels at understanding tasks:

Classification Tasks:

Sentiment Analysis
Topic Classification
Spam Detection
Intent Recognition

Token-Level Tasks:

Named Entity Recognition (NER)
Part-of-Speech Tagging
Word Sense Disambiguation

Span-Level Tasks:

Question Answering
Extractive Summarization
Information Extraction

Sentence-Pair Tasks:

Semantic Similarity
Natural Language Inference
Paraphrase Detection

Industry Impact

BERT powers:

- Google Search (understanding queries)
- Customer service chatbots
- Content moderation
- Document understanding

Case Study: Google Search

BERT revolutionized search in 2019

Problem:

Search queries often depend on context and word order
Traditional keyword matching misses nuance
Example: "2019 brazil traveler to usa need a visa"

Before BERT:


Focused on "brazil" and "usa"
Missed importance of "to"
Returned results about US travelers to
Brazil
Wrong direction!



Wrong: US → Brazil

With BERT:

Understands "to" is critical
Grasps full context and direction
Returns correct results about Brazilian
travelers to US
Correct understanding!



Correct: Brazil → US

Google reported BERT improved 1 in 10 searches in English

Question Answering with BERT

Extractive QA: Find answer span in passage

Example

Context: "The Normans (Norman: Nourmands; French: Normands; Latin: Normanni) were the people who in the 10th and 11th centuries gave their name to Normandy, a region in France."

Question: "In what country is Normandy located?"

Answer: France

```
from transformers import pipeline

# Load QA pipeline with BERT
qa_pipeline = pipeline("question-answering", model="bert-large-uncased-whole-word-masking-finetuned-squad")
```

Named Entity Recognition

Token-level classification task

Example

Input: "Apple Inc. is headquartered in Cupertino, California."

Output:

- Apple Inc. □ ORGANIZATION
- Cupertino □ LOCATION
- California □ LOCATION

```
from transformers import pipeline

# Load NER pipeline
ner_pipeline = pipeline("ner", model="dslim/bert-base-NER")
```

```
# Extract entities
```

Sentiment Analysis 😊 😐 😞

Sequence classification task

Examples

- "This movie was absolutely amazing!" → POSITIVE
- "The product broke after one week." → NEGATIVE
- "The weather is cloudy today." → NEUTRAL

```
from transformers import pipeline

# Load sentiment analysis pipeline
sentiment_pipeline = pipeline("sentiment-analysis", model="
    distilbert-base-uncased-finetuned-sst-2-english")

# Analyze sentiments
texts = [
```

Semantic Similarity

Measuring sentence similarity with BERT embeddings

```
from transformers import BertTokenizer, BertModel
import torch
import torch.nn.functional as F

tokenizer = BertTokenizer.from_pretrained('bert-base-uncased')
model = BertModel.from_pretrained('bert-base-uncased')

def get_sentence_embedding(sentence):
    inputs = tokenizer(sentence, return_tensors='pt', padding=True
                        , truncation=True)
    outputs = model(**inputs)
    # Use [CLS] token embedding as sentence representation
    return outputs.last_hidden_state[:, 0, :]
```


Cognitive Neuroscience Perspective

How do brains and models process language?

Predictive Processing in the Brain:

- Brain constantly predicts upcoming input
- N400: Neural response to unexpected words
- P600: Syntactic anomaly detection
- Context shapes predictions
- Prediction errors drive learning

Key Brain Regions:

- Left IFG:** Syntax processing
- Left STG/MTG:** Semantic processing
- ATL:** Conceptual knowledge

Predictive Processing in Models:

- BERT: Predict masked words
- GPT: Predict next word
- Both use context to predict
- Surprise = high loss
- Gradient descent = learning

Similarities:

- Both hierarchical
- Both context-sensitive
- Both predictive
- Both learn from errors

References: Kuperberg & Jaeger (2016), Willems et al. (2016), Hagoort & Indefrey (2014)

Prediction in Brains vs. Language Models

Parallels between neural and artificial systems

Human Brain	Transformer Models
N400 component (surprise at unexpected word) (low probability prediction)	High cross-entropy loss
Hierarchical processing (sounds → words → sentences) (tokens → phrases → meaning)	Layer-by-layer representations
Context integration (prior discourse, world knowledge) (attend to relevant context)	Self-attention mechanism

Neural Encoding with Language Models

Can we predict brain activity from language models?

Approach:

- Show participants text while recording brain activity (fMRI/EEG)
- Extract representations from language model (e.g., BERT layer 6)
- Train regression model: Model embeddings $\rightarrow$ Brain activity
- Test: Can we predict brain responses to new text?

Findings:

- **Better models $\rightarrow$ Better brain prediction**
- BERT outperforms Word2Vec at predicting brain activity
- Different layers correlate with different brain regions
- Middle layers best for semantic areas
- Suggests shared representations?

But...

Correlation $\neq$ Causation! Models might capture statistical regularities that brains use, without using the same mechanisms.

Reference: Willems et al. (2016) - "Prediction during natural language comprehension"

Discussion: What Does the Model "Understand"? 🤔

Does BERT understand language?

Evidence FOR understanding:

- Captures syntax and semantics
- Resolves ambiguity
- Handles long-range dependencies
- Generalizes to new examples
- Predicts brain activity
- Solves complex tasks

"If it acts like it understands, maybe it does?"

Evidence AGAINST understanding:

- No grounding in physical world
- No sensory experience
- No social context
- Brittle to adversarial examples
- No common sense reasoning
- Only pattern matching?

"Understanding requires more than statistical patterns"

Key Questions

- What is the difference between understanding and correlation?
- Can meaning exist without grounding?
- Is human understanding fundamentally different?

Class Discussion: What do YOU think?

Adversarial Examples and Brittleness ⚠

BERT can be fooled easily

Example: Sentiment Analysis

Original: "This movie was great!"

Prediction: POSITIVE ☐

Modified: "This movie was great! [SEP] terrible terrible terrible terrible"

Prediction: NEGATIVE ☐

Just adding repeated negative words after [SEP] flips the prediction!

Other Brittleness Issues

- **Word substitutions:** Replace "good" with synonym "excellent" ☐ wrong prediction
- **Paraphrases:** Semantically equivalent but different words ☐ different predictions

Limitations of Current Models ⚠️

Despite impressive performance, transformers have limitations:

Quadratic Complexity

- Self-attention scales as $O(n^2)$
- Limited context windows (512-4096 tokens)
- Cannot process very long documents efficiently

No True Understanding

- Pattern matching vs. comprehension
- Lack of common sense
- No world model

Data Efficiency

- Requires massive training data
- Humans learn language with much less data
- Not biologically plausible

Biases and Fairness

- Inherits biases from training data
- Can amplify stereotypes
- Ethical concerns

Bias in Language Models

Models reflect and can amplify societal biases

Examples of Bias

- **Gender bias:**
 - "The doctor said [MASK] was running late" → he (90%)
 - "The nurse said [MASK] was running late" → she (80%)
- **Occupational stereotypes:**
 - Associates certain jobs with certain genders
 - "Programmer" → male pronouns
 - "Secretary" → female pronouns
- **Racial bias:**
 - Different sentiment for names associated with different races
 - Can perpetuate harmful stereotypes

Mitigation Strategies:

Debiasing techniques during training
Balanced and diverse training data
Post-processing and filtering
Human oversight and auditing
Ongoing research area!

Practical Tips for Working with Transformers 💡

Start with Pre-trained Models

- Don't train from scratch (too expensive!)
- Use HuggingFace Model Hub
- Choose appropriate model size

Fine-tuning Best Practices

- Use small learning rate ($1e-5$ to $5e-5$)
- Add warmup steps
- Monitor for overfitting
- Freeze early layers if data is limited

Computational Efficiency

- Use mixed precision training (FP16)
- Gradient accumulation for larger batch sizes
- Consider DistilBERT for faster inference
- Use FlashAttention when available

Evaluation

- Use task-specific metrics
- Test on out-of-distribution data
- Check for biases
- Visualize attention for interpretability

Deployment Considerations

Moving from research to production

Model Selection

- Balance quality vs. speed
- DistilBERT for low-latency applications
- BERT-Base for balanced performance
- BERT-Large when quality is critical

Optimization

- Quantization: FP32 $\rightarrow$ INT8 (4x smaller, faster)
- Model pruning: Remove unnecessary weights
- Knowledge distillation: Compress to smaller model
- ONNX Runtime for optimized inference

Infrastructure

- CPU acceleration for batch processing
- CPU optimization for single requests
- Caching for repeated queries
- Load balancing for high traffic

Monitoring

- Track latency and throughput
- Monitor prediction quality
- Detect distribution drift
- A/B testing for improvements

Future Directions

Where is the field heading?

Longer Context

- Efficient attention mechanisms (linear, sparse)
- Models with 100K+ token context
- Better long-document understanding

Multimodal Models

- Vision + Language (CLIP, DALL-E)
- Audio + Language (Whisper)
- Grounded understanding

Better Pre-training

- More efficient objectives
- Curriculum learning
- Continual learning

Smaller, More Efficient Models

- Better compression techniques
- Lottery ticket hypothesis
- Edge deployment

Addressing Limitations

- Debiasing and fairness
- Robustness and adversarial training
- Common sense reasoning
- Interpretability and explainability

Encoder vs Decoder Models Revisited

Different models for different tasks

	Encoder (BERT)	Decoder (GPT)
Architecture	Bidirectional self-attention	Unidirectional (causal) attention
Pre-training	Masked Language Modeling	Next token prediction
Best for	Understanding tasks:	Generation tasks:
• Classification • NER, QA • Similarity • Text completion • Dialogue • Creative writing		

Discussion Questions

Understanding vs. Pattern Matching:

Where do you draw the line?

Is there a test for "true" understanding?

Does it matter for applications?

Brain-Model Parallels:

How useful are these comparisons?

What can neuroscience learn from AI?

What can AI learn from neuroscience?

Bias and Fairness:

Who is responsible for addressing bias?

Can we ever have completely unbiased models?

How do we balance accuracy and fairness?

Future of NLP:

Will encoder models remain relevant?

Are decoder-only models the future?

What's the next big breakthrough?

Assignment 4: Context-Aware Models

Hands-on experience with transformers!

Tasks:

Implement Attention Mechanism

- Build scaled dot-product attention from scratch
- Visualize attention weights

Fine-tune BERT

- Load pre-trained BERT
- Fine-tune on sentiment analysis
- Compare to baseline models

Analyze Contextual Embeddings

- Extract embeddings for polysemous words
- Visualize how context changes representations
- Compare BERT vs Word2Vec

Explore Different Architectures

- Compare BERT (encoder) vs GPT (decoder)
- Test on different tasks
- Analyze strengths/weaknesses

Research Component

- Read one paper from references
- Write brief summary & critical analysis

Due: Check course website for deadline

Summary: Weeks 5-6

What we learned:

Evolution of Context

- Seq2Seq → Attention → Transformers
- From bottleneck to full parallelization

Transformer Architecture

- Self-attention, multi-head attention
- Positional encoding, layer norm, residuals
- Encoder-only (BERT), Decoder-only (GPT), Both (T5)

BERT & Variants

- Masked Language Modeling
- Pre-train then fine-tune paradigm
- RoBERTa, ALBERT, DistilBERT, ELECTRA

Applications

- Classification, NER, QA, similarity
- Real-world impact (Google Search, etc.)

Broader Implications

- Brain-model parallels
- Understanding vs. pattern matching
- Limitations and future directions

Resources & Further Reading

Key Papers:

Vaswani et al. (2017) - Attention Is All You Need

Devlin et al. (2019) - BERT: Pre-training of Deep Bidirectional Transformers

Liu et al. (2019) - RoBERTa

Sanh et al. (2019) - DistilBERT

Dao et al. (2022) - FlashAttention

Cognitive Neuroscience:

Hagoort & Indefrey (2014) - The neurobiology of language beyond single words

Kuperberg & Jaeger (2016) - What do we mean by prediction in language comprehension?

Willems et al. (2016) - Prediction during natural language comprehension

Tutorials:

HuggingFace Course: Chapters 1, 7

The Illustrated Transformer:

<https://jalammar.github.io/illustrated-transformer/>

BertViz: <https://github.com/jessevig/bertviz>

Looking Forward in the Course

Where do we go from here?

Upcoming Topics:

Week 7: Decoder models and text generation (GPT family)

Week 8: Scaling laws and large language models

Week 9: Prompting, in-context learning, and instruction tuning

Week 10: Alignment, RLHF, and ethical considerations

The Journey Continues:

From understanding (BERT) to generation (GPT)

From supervised learning to few-shot learning

From narrow tasks to general-purpose models

From academic research to societal impact

The transformer revolution continues! 

Questions? 

Discussion Time

Topics to discuss:

- BERT applications
- Understanding vs. pattern matching
- Cognitive neuroscience connections
- Limitations and future work
- Assignment 4 questions

Thank you! 

See you in Week 7 for GPT and text generation!